"AI-Powered Study Assistant using Retrieval-Augmented Generation (RAG)"

Chalini M – 2023BCSE07AED440
Spandana Basava Kulal – 2023BCSE07AEd441
Devika A S – 2023BCSE07AEd495
G Jyostna – 2023BCSE07AED518

A CASE STUDY REPORT

Submitted in partial fulfilment of the requirements for the

Course Name: LLM
Course code: CS6012

BACHELOR OF TECHNOLOGY

Branch: CSE Specialization: AIML-E

Of Alliance University

April 2026

ALLIANCE SCHOOL OF ADVANCED COMPUTING

ALLIANCE UNIVERSITY, BENGALURU

CONTENTS

S.No	Title	Page No
	ABSTARCT	
1.	INTRODUCTION	4
2.	PROBLEM STATEMENT	5
3.	CORE CONCEPT - RAG	6
4.	SYSTEM ARCHITECTURE	7
5.	WORKING FLOW	8
6.	MODULES OF THE SYSTEM	9
7.	TECHNOLOGIES USED	10
8.	EXAMPLE SCENARION	11
9.	ADVANTAGES AND LIMITATIONS	12
10.	CONCLUSION AND FUTURE SCOPE	13
11.	REFERENCES	14
12.	PLAGARISM REPORT	15

ABSTRACT

This case study focuses on developing an AI-powered study assistant using Retrieval-Augmented Generation (RAG). Many students find it difficult to search through large amounts of study material such as PDFs, notes and textbooks. At the same time, traditional language models may sometimes give incorrect or general answers because they rely only on pre-trained knowledge.

To overcome this issue, the proposed system combines information retrieval with answer generation. It first searches for relevant content from the uploaded study materials answers. This helps in providing responses that are based on actual content rather than assumptions. The system is built using python along with tools like LangChain, and GPT- 4. Overall, this study show how RAG can be used to improve learning by making information access faster, easier, and more reliable for students.

INTRODUCTION

Large Language Models (LLMs) are powerful AI systems that can understand and generate human-like text. They are widely used in applications such as chatbots, content writing, and educational tools. In recent years, students have increasingly depended on digital study materials like PDFs, lecture notes and online resources for learning. However, finding specific information from these large documents is often difficult and time-consuming. Students usually need to go through multiple files to get the required answer, which reduces efficiency and wastes time.

At the same time, traditional AI models may provide general or sometimes incorrect answers because they depend only on pre-trained knowledge and do not have access to specific study materials. This creates a need for a smarter system that can understand user queries and provide accurate answers directly from the given documents. This project focuses on developing an AI-based study assistant that helps students quickly find relevant information from their study materials and improves their learning experience.

PROBLEM STATEMENT

Students often deal with large amounts of study material such as PDFs, lecture notes and textbooks which are not always well organized. Finding a specific answer from these materials takes a lot of time and effort. In many cases, important information is spread across different documents, making it difficult for students to quickly understand a concept.

Manual searching is not only slow but also inefficient, especially during exams or revisions. Existing tools like basic search functions or simple chatbots do not always give accurate or relevant answers because they lack proper understanding of the content. As a result, students may struggle to find the right information at the right time, which can affect their learning efficiency and overall performance.

CORE CONCEPT – RAG

Retrieval-Augmented Generation (RAG) is an approach that improves how AI systems generate answers by combining two processes retrieving information and generating responses. Instead of depending only on the knowledge the model learned during training, the system first searches for relevant information from available documents such as PDFs or notes.

Once the relevant content is found, it is given as input to the language model, which then generates a clear and meaningful answer based on that information. This makes the response more accurate and related to the actual content rather than general knowledge. By using this method, the chances of incorrect or misleading answers are reduced, and the system becomes more reliable for tasks like answering questions from study materials.

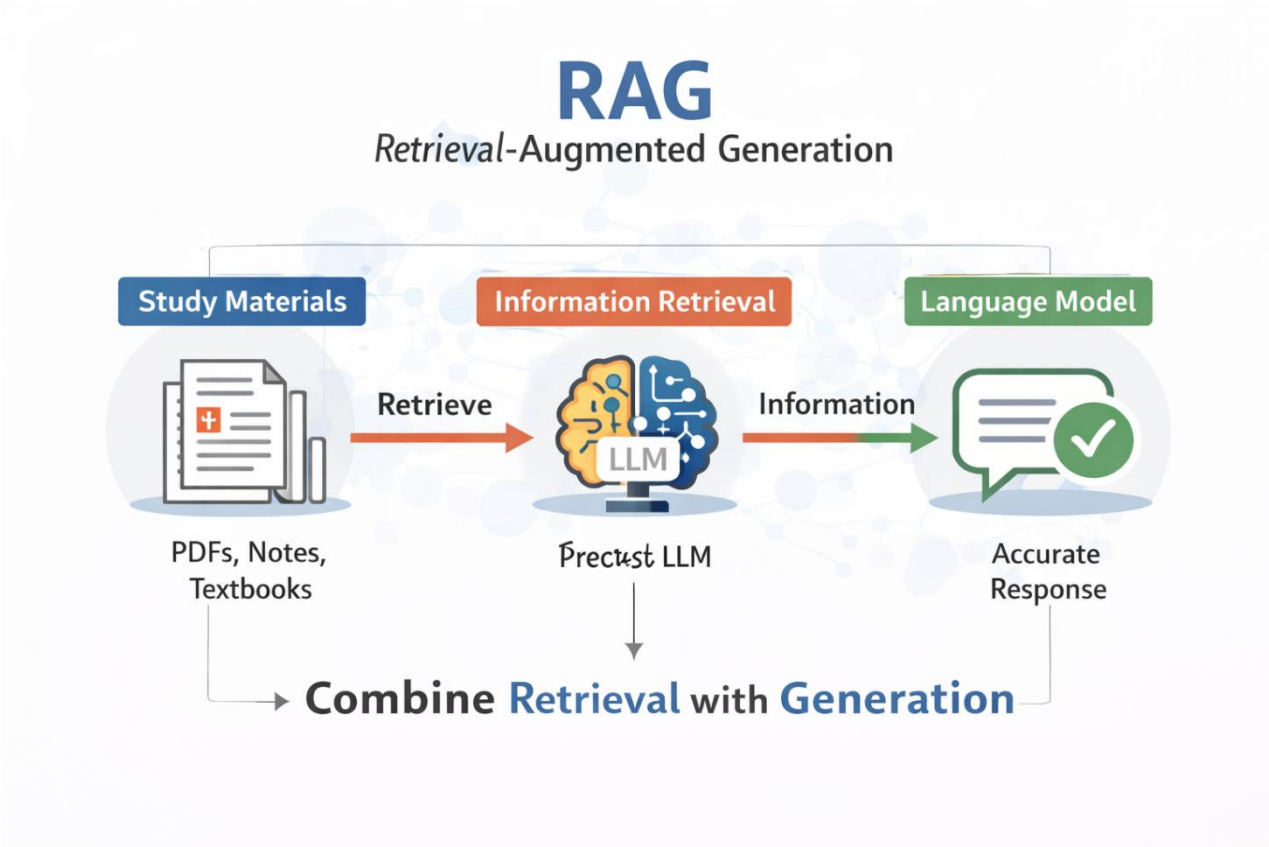

SYSTEM ARCHITECTURE

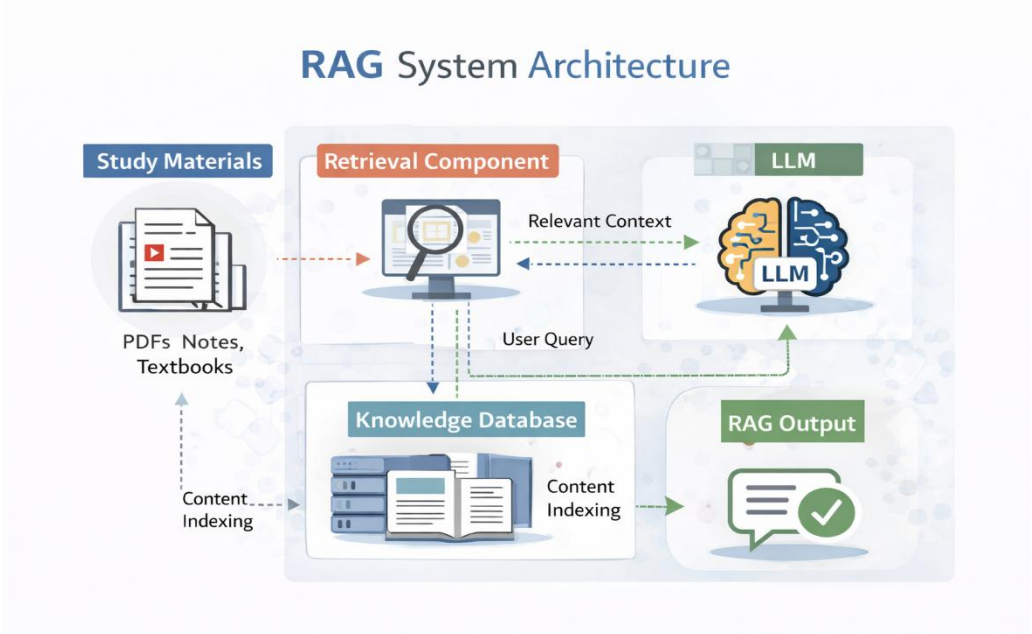

The system architecture of the AI-powered study assistant is designed to process user queries and generate accurate answers using a structured approach. It consists of different components that work together to retrieve and generate information effectively. When a user asks a question, the system first process the query and sends it to the retriever. The retriever searches for relevant information from the stored documents in the vector database. These documents are already converted into embedding's and stored for efficient searching.

After retrieving the most relevant content, the information is passed to the language model. The mode then analyzes both the user query and the system provides the generated answers to the user. This architecture ensures that the answers are not only generated by the model but are also based on actual data, making the system more reliable and useful for students.

WORKING FLOW

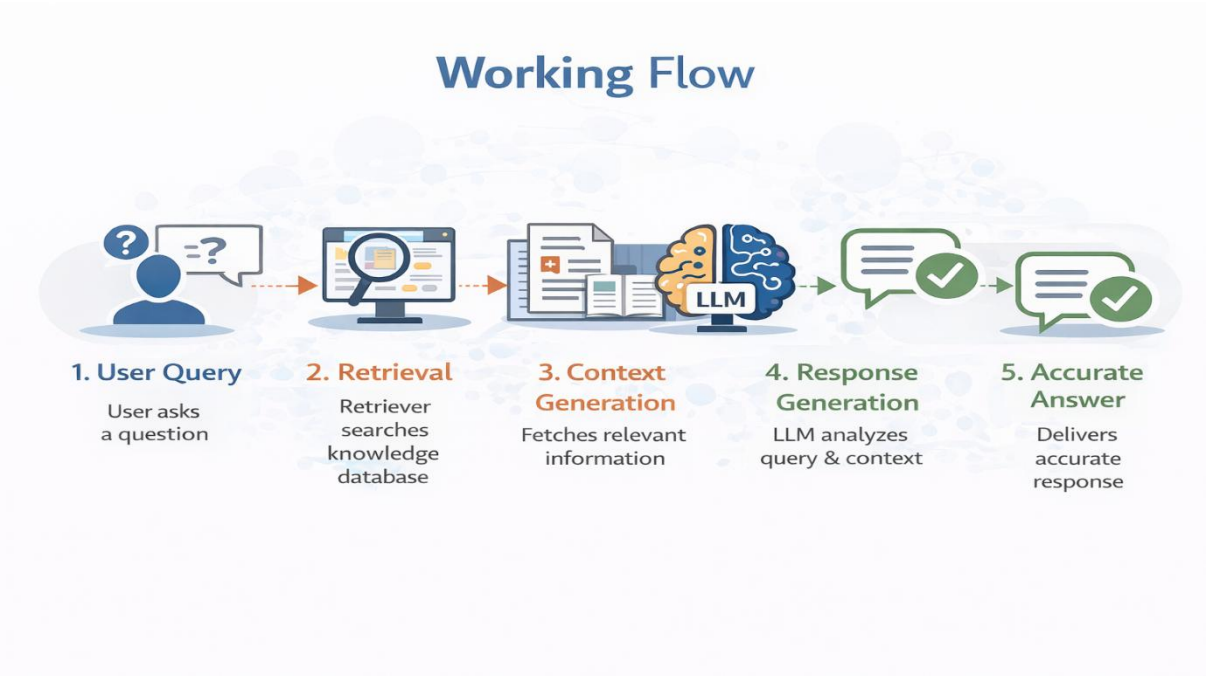

The working flow of the AI-powered study assistant begins when the user uploads study materials such as PDFs or notes into the system. These documents are first processed and divided into smaller parts to make them easier to handle. Each part is then converted into numerical representations called embedding's, which help the system understand the meaning of the text. These embedding's are stored in a vector database for quick and efficient searching

When a user asks a question, the system converts the query into an embedding and searches the database to find the most relevant information. The retrieved content is then passed to the language model along with the users questions. Based on this information, the model generates a clear and accurate answer. Finally, the system displays the answer to the user, ensuring that it is based on actual study material.

MODULES OF THE SYSTEM

The AI-powered study assistant is divide into several modules, each responsible for a specific task in the system. These modules work together to process the documents and generate accurate answers for user queries.

The first module is the document loader, which reads and extracts text from uploaded files such as PDFs and notes. After that, the next chunking module splits the large text into smaller parts so that it can be processed efficiently.

Next, the embedding module converts these text chunks into numerical vectors, which help the system understand the meaning of the content. These vectors are then stored in the vector database, which allows quick searching of relevant information.

The retriever module plays an important role by finding the most relevant data based on the users query. Finally, the language model processes the retrieved information along with the query and generates the final answer. Together, these modules ensure that the system works smoothly and provides accurate, context-based responses to users.

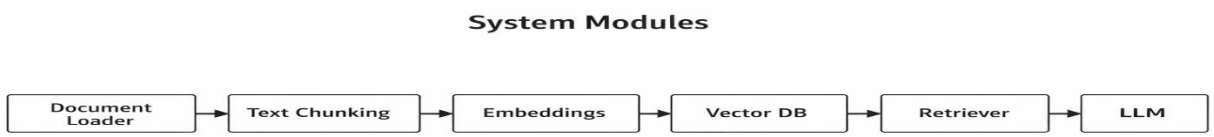

TECHNOLOGIES USED

1. **Python** – is used as the main programming languages for developing the system. It is simple and widely used for Ai application

2. **FASIS** – is used as a vector database to store embedding's and perform fast similarity search for retrieving relevant data

3. **LangChain** – is used to manage the workflow and connect different components such as document processing, retrieval and generation

4. **GPt-4** – is the language model used to generate accurate and meaningful answers based on the retrieved information

EXAMPLE SCENARIO

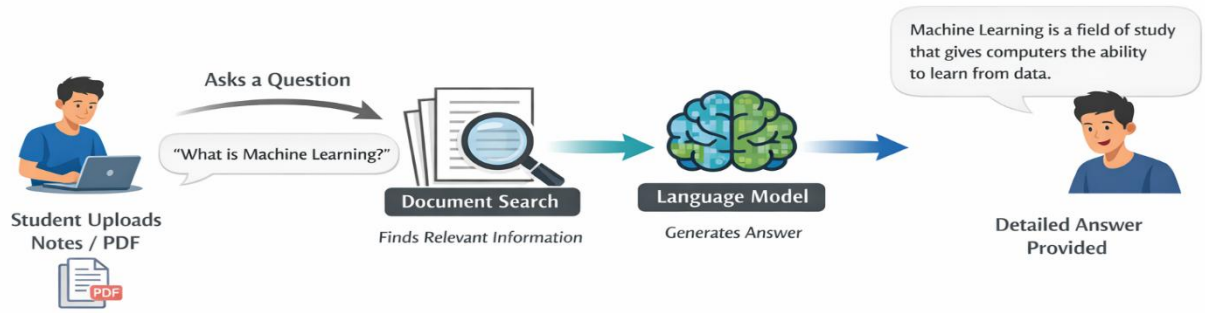

To understand how the system works, consider a simple example. A student uploads study materials such as notes or PDF related to Machine Learning. The student then asks a question like "What is Machine Learning"?

The system processes this query by searching for relevant information from the uploaded documents. It retrieves the most useful content and passes it to the language model. Based on this information, the model generates a clear and accurate answer. This ensures that the response is not general but directly based on the students study material making it more useful and reliable.

ADVANTAGES

1. Provides accurate answers based on actual study material
2. Reduces incorrect or misleading response
3. Saves time by avoiding manual searching
4. Helps students learn more efficiently
5. Works with user-specific data

LIMITATIONS

1. Depends on the quality of uploaded documents
2. May take slightly more time due to retrieval process
3. Requires initial setup of documents
4. Performance depends on embedding quality

CONCLUSION

The RAG – Based study assistant improves the process of finding information from study materials by combining retrieval and generation. It provides accurate and context-based answers, which helps students save time and learn more effectively. This project shows how AI can be used education to make learning easier and more efficient.

FUTURE SCOPE

1. Voice-based interaction
2. Multi-language support
3. Mobile application development
4. Integration with learning platforms
5. Improved accuracy with advanced models

REFERENCES

1. Lewis, P., Perez, E., Piktus, A., et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. Advances in Neural Information Processing Systems.
2. Gao, Y., Xiong, Y., Gao, X., et al. (2024). Retrieval-Augmented Generation for Large Language Models: A Survey. arXiv.
3. Johnson, J., Douze, M., & Jégou, H. (2019). Billion-Scale Similarity Search with GPUs. IEEE Transactions on Big Data.
4. OpenAI. (2023). GPT-4 Technical Report. arXiv.
5. Vaswani, A., Shazeer, N., Parmar, N., et al. (2017). Attention Is All You Need. Advances in Neural Information Processing Systems.
6. Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT Networks. EMNLP.
7. Chase, H. (2022). LangChain Documentation. GitHub Repository.
8. Facebook AI Research. (2023). FAISS Documentation.

PLAGARISM REPORT